

Assignment A1: Team 41

Sven Steens

Riccardo Torlaini

Rafał Polak

21/11/2022

1 Introduction

Note that due to ambiguity within the assignment/rules description, we will be slightly altering the naming. We will be calling:

- A vertical set of numbers - a **column**
- A horizontal set of numbers - a **row**
- A $n \times m$ rectangular set of numbers - a **region**
- Any of the above generically - a **group**
- A "cell" which can be filled with precisely one number - a **square**

Such naming is used both by this report and inside the code. Additionally, some functions in the code may use the term *amount* to refer to cardinalities of countable sets, but it is not to be confusing with the term *number* which we frequently use to refer to the digit value being placed on a sudoku grid.

All the main ideas and ways the agent operates should be understandable from the report and the well-commented code, but in case the reviewer would like to get more explanation, refer to Section 5 (Code overview).

2 Agent Description

The main functionality of the agent is - obviously - the minimax algorithm to evaluate the (current and possible) board states. The other notable technique employed by this agent is the *alpha-beta pruning* that significantly sped up the algorithm in our tests (around 6 times faster computations for the random-2x3 board). Since the filled number does not matter for the score, it does not matter for the minimax. As a result, the algorithm works in two steps. Firstly we use a custom sudoku solver to determine a feasible number for each square. Then we find the optimal square to fill in. This optimal square is searched with respect to some depth on the minimax search (typically starting with depth=1, i.e. considering the game states only 1 move ahead of the current board state) and after such depth was thoroughly explored, move to a higher depth. This procedure carries on until some maximum depth for the current iteration. After each finalized such exploration, the best move (according to the current depth) is reported, with a new one overwriting the previous "best-declared move" after the deeper search has been conducted. Afterwards, increase the maximum depth and re-iterate. The maximum depth parameter is in practice unbounded - given infinite time, the entire game tree will be traversed. The traversal uses a depth-first search at each iteration. Contrary to suspicion, saving the current game tree when moving to a higher depth, did not prove significant in speeding up the algorithm, given the amount of time needed to save and load the time, as opposed to how quick our implementation naturally is.

2.1 Heuristics

We chose not to use any user-defined heuristics for evaluating the positions, other than the best score (assuming perfect play from the opponent) resulting in playing such a move (up to some depth). We noticed how the agent restrains from filling up squares that have an even number of empty squares in their corresponding groups. That makes intuitive sense - after simultaneous updates to the same groups, the opposing player will put the last number in and receive the points for these groups, so that is not our dream scenario.

2.2 "Solve_sudoku" function

Our sudoku solver finds a solution for a given Sudoku board in a reasonably "human" fashion: first, it attempts to iteratively fill in all the empty squares that can be simply determined to have a correct solution - for they are the only empty squares in their corresponding group. After any such square is filled, the procedure is restarted. Afterwards, the function makes a random guess, filling in some other empty square. Then, it proceeds with the previous step. These steps are invoked iteratively, until either the sudoku is solved, or turns out to be unsolvable. In the latter case, The most recent guess has to be undone and another one is made. As soon as no other possible guesses at this board state remain, the algorithm goes another step back and undoes the guess before. This can theoretically keep happening until the very first guess of the board was made. Further optimizations can be made to that function, however as it is only being called once per game, its running time is not such impactful of a bottleneck.

2.3 Alpha-beta Pruning

Of course, we can not explain the alpha-beta pruning better than this video does: <https://youtu.be/l-hh51necgDI>. But in case you prefer to read instead consider the following example where a positive score favours Player1 ("White") and a negative score favours Player2 ("Black"): Imagine at some point in the game tree, White can make a move that leads to an evaluation score of 1.0 with Black's perfect play. The other move (we assume 2 moves only) leads to Black's choice between a move evaluated at -1.5 and another move branch. Then, this branch does not have to be explored at all - Black will choose a move that gives them at least 1.5 point advantage, so regardless of specifics of this branch, White should not choose that move; hence, we can "*prune*" that branch and terminate this part of minimax, evaluating the position as 1.0 at White-to-move turn.

2.4 Technical Notes

A few techniques were used to reduce the running time of each minimax iteration. Instead of recalculating the possible moves (`open_squares` in the code) at each new board state, the possible moves are calculated at the start. Whenever a move is done, it is then removed from this list before it is passed on to the child. To add to this, we also do not fully recalculate how many points a move can earn at each state. Instead, we keep track of the amount of empty (zero) squares in each row, column and region (`empty_squares` in the code). Thus we know that if one of the three groups of a square has exactly 1 still empty square, placing that square will complete that group, thus earning points. After a move is done we can then decrease the number of empty squares in the groups of that move by one. Lastly, we avoid copying lists and dictionaries as much as possible. Copying these structures is costly if we were to do it in every node. Instead, we change these variables before going to a lower depth and change them back after we are done. This is especially more time efficient in the case of the `open_squares` variable, where we remove and then read one value, as opposed to copying a whole list.

3 Agent Analysis

3.1 Formal results on 3x3 boards

The **Results** are formatted in a way that 'w' means a victory (win) of our agent, 'l' its loss and 'd' a draw. See Table 1 below for the detailed scores.

Table 1: Summary of the scores when testing the agent against the provided players.

Opponent	Board	Timer	Results	Winrate
random_player	random-3x3	0.5s	wwwwwwwwww	100%
random_player	random-3x3	1s	wwwwwwwwww	100%
random_player	random-3x3	5s	wwwwwwwwww	100%
greedy_player	random-3x3	0.5s	llwlwwlwl	50%
greedy_player	random-3x3	1s	llwwlllll	20%
greedy_player	random-3x3	5s	wwlwlwwl	60%
random_player	empty-3x3	0.5s	wllwllwl	30%
random_player	empty-3x3	1s	llllwllll	10%
random_player	empty-3x3	5s	wlwlwwlwwl	60%
greedy_player	empty-3x3	0.5s	llllwvlll	20%
greedy_player	empty-3x3	1s	llwwlllll	20%
greedy_player	empty-3x3	5s	wwwwlwlww	80%

To better visualize these results and give more insight into their significance, we have also plotted the scores for each player and board combination at all three timer settings:

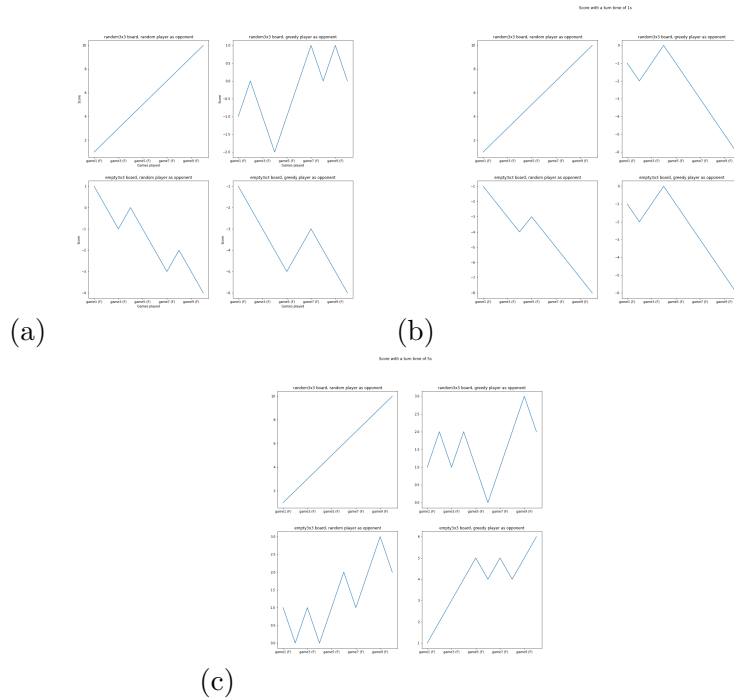


Figure 1: (a) Timer = 0.5s (b) Timer = 1s (c) Timer = 5s

3.2 For smaller boards

For the empty-2x2 board, a simple pattern emerged. At any timer and against any opponent, the player moving 2nd, showcased a much higher win-rate than any agent who took the

first move. It makes intuitive sense! With an even number of squares to be filled in, the player that makes the 2nd move will also be the one making the last move of the game - and *this* is truly "overpowered"! Making the last move is *guaranteed* to bring the player additional 7 points, for the last move is forced to complete all 3 groups - a column, a row and a region. For a game on a large board, these guaranteed 7 points will still play a significant role, but it is less of a critical factor. For all the smaller boards than 3x3 - the impact of the last move will be tremendous. Notice how this problem also is applicable to boards with an even number of squares to be filled - then, it will be the "going first" player who gets the advantage.

4 Reflection

By far, the strongest point of our agent is how exhaustive the search is. Given a long timer per turn, it would theoretically be able to fully explore the game tree and "solve" the current competitive Sudoku board. Furthermore, the included alpha-beta pruning as well as optimizations in how the minimax is run result in searches over 15x faster than a 'naive' implementation of the algorithm.

There are however some limitations to the algorithm. First of all, the elephant in the room, the performance is sub-optimal. Due to the theoretical groundedness of the algorithm and the positive results of previous practical tests we believe this is due to an uncaught bug in the implementation. At the time of writing it is however unfortunately too late to fix this. See this as an exercise for the reader. On more theoretical notes, due to the fact that the amount of 'thinking' time is unknown, the algorithm essentially has to do many different runs of minimax, instead of a single deep one, costing valuable time. Additionally, while not allowed by the rules, an optimal version of this algorithm would send information between moves. Specifically, it would store the results of a tree of depth n , so that during its next move it would not have to calculate minimax till depth $n-2$, but could use the already established calculations instead. Further optimizations can of course also be implemented, such as the inclusion of transposition tables and improvements to our sub-optimal Sudoku solver. Lastly, our algorithm is missing the tactical ability to intentionally skip moves. As per the rules, the AI could intentionally pick a legal, non-taboo move that makes the Sudoku unsolvable in order to skip their turn without filling in any squares. This proved too difficult to implement into our framework but will put the AI at a disadvantage against AIs that are capable of this.

5 Code overview

Below, we would like to give a very top-level overview of the code appended, so that the reviewers can understand it more easily.

Lines 21-23: At initialization, we retrieve the info on the open (available) squares on the board, how many open squares are there in each group and what numbers are missing in each group.

Line 26: We solve the sudoku using the approach given in 2.2.

Lines 29-32: We run the minimax algorithm on the current game state to retrieve (and propose) the best move in the current situation. The minimax will be running recursively, each time starting with depth=1 and increasing it until max_depth, where max_depth will increase by 1 at every iteration.

Lines 34-48: We simply update the "proposed move" using this function.

Lines 88-91: Calculating how many point would the given move contribute to the overall result (bearing in mind that this will be a negative value if it is the turn of the minimizing player, as determined in Line 81).

Lines 93-97: Since this current move is being considered, the numbers of empty groups have to be updated accordingly, if we decide to play it, so that is what we do here.

Lines 99-100: Recurse minimax with a larger ("going more down the tree") depth. This way, the changes to the numbers of empty groups and game score, altered by considering the current move, stay in place and have an impact on this entire "branch" of the game tree that we will be traversing by recursing minimax in this place.

Lines 102-111: Since in this place we finished inspecting the "previous" branch of the game tree, we restore the values of numbers of empty squares per group. We also save the score of the previous branch in case this was the best one seen so far (so that we might want to choose the move leading to this branch in the game).

Lines 113-117: See Section 2.3 (Alpha-beta pruning).

Lines 121-133: A function to retrieve all the empty squares (where inputting a number is possible) from the current game state.

Lines 135-170: The actual function to retrieve the numbers of empty squares per group, as invoked in Lines 21-23.

Lines 172-204: The actual function to retrieve the missing numbers per group, as invoked in Lines 21-23.

Lines 206-260: Very naïve sudoku solver, explained in Section 2.2 ("Solve_sudoku" function).

Python files

Code Listing 1: sudokuai.py.

```
1 from copy import deepcopy
2 from competitive_sudoku.sudoku import GameState, Move, SudokuBoard

4 class SudokuAI(object):
5     """
6     Sudoku AI that computes a move for a given sudoku configuration .
7     """
8     def __init__(self):
9         self.best_move = [0, 0, 0]
10        self.lock = None

12    def compute_best_move(self, game_state: GameState) -> None:
13        """
14        The AI calculates the best move for the given game state .
15        It continuously updates the best_move value until forced to stop .
16        Firstly , it creates a solved version of the sudoku such that it has a valid number for every square .
17        Then it repeatedly uses minimax to determine the best square , at increasing depths .
18        @param game_state: The starting game state .
19        """
20        # Calculates the starting variable minimax needs
21        open_squares = game_state.board.get_open_squares()
22        empty_squares = game_state.board.get_empty_squares()
23        numbers_left = game_state.board.get_numbers_left()

25        # Gives a solution of the board
26        solved_board = solve_sudoku(deepcopy(game_state.board), deepcopy(open_squares), numbers_left)

28        # Calculate for every increasing depth
29        for depth in range(1,9999):
30            move = minimax(max_depth = depth, open_squares = open_squares, empty_squares = empty_squares, m =
31                               game_state.board.m, n = game_state.board.n)[1]
32            number_to_use = solved_board.get(move[0], move[1])
33            self.propose_move(Move(move[0], move[1], number_to_use))

34    def propose_move(self, move: Move) -> None:
35        """
36        Updates the best move that has been found so far .
37        N.B. DO NOT CHANGE THIS FUNCTION!
38        This function is implemented here to save time with importing
39        @param move: A move.
40        """
41        i, j, value = move.i, move.j, move.value
42        if self.lock:
43            self.lock.acquire()
44            self.best_move[0] = i
45            self.best_move[1] = j
46            self.best_move[2] = value
47        if self.lock:
48            self.lock.release()

50    #The below functions exist so that we can create certain references in the minimax function
51    def greater(i: int, j: int) -> int:
52        return i > j

54    def smaller(i: int, j: int) -> int:
55        return i < j

57    def minimax(max_depth: int, open_squares: list, empty_squares: dict, m: int, n: int,
58    is_maximizing_player: bool = True, current_score: int = 0, alpha: int = float("-inf"), beta: int = float("inf")):
59        """
60        A version of the minimax algorithm implementing alpha-beta pruning .
61        Every time we create a child , we calculate how many points the move associated with that child might get us .
62        This calculation is done with empty_squares, while all potential moves are kept track of via open_squares .
63        Variables with default values take those values during the first iteration .
64        @param max_depth: The maximum depth the function is allowed to further search from its current depth .
65        @param open_squares: A list containing all still open squares/possible moves .
66        @param empty_squares: A dictionary containing the number of empty squares for each group .
67        @param m: The number of rows per region for this board, used to calculate regions from coordinates .
68        @param n: The number of columns per region for this board, used to calculate regions from coordinates .
69        @param is_maximizing_player: Whether the current player is attempting to maximize or minimize the score .
70        @param current_score: The score at the node we start this iteration of minimax on .
71        @alpha: The alpha for alpha-beta pruning .
72        @beta: The beta for alpha-beta pruning .
73        @return: The score that will be reached from this node a maximum depth and the optimal next move to achieve that .
74        """

76        # If we have hit either the maximum depth or if there are no more moves left , we stop iteration
77        if max_depth == 0 or not open_squares:
78            return current_score, (-1,-1)

80        # Switches values around depending on if the player is maximizing or not
81        value, function, multiplier = (float("-inf"), greater, 1) if is_maximizing_player else (float("inf"), smaller, -1)

83        # Initializes where we store the best move and its associated score of this node
84        best_score = value
85        best_move = open_squares[0]

87        for move in open_squares[:]:

89            # Calculates how the move would change the score
90            amount_finished = (empty_squares["row"][move[0]] == 1) + (empty_squares["column"][move[1]] == 1) + (empty_squares
91                ["region"][int(move[0] / m)*m + int(move[1] / n)] == 1)
92            new_score = current_score + multiplier*{0:0, 1:1, 2:3, 3:7}[amount_finished]
```

```

93     # Removes the move from open_squares and updates empty_squares to account for the move
94     open_squares.remove(move)
95     empty_squares["row"][move[0]] -= 1
96     empty_squares["column"][move[1]] -= 1
97     empty_squares["region"][int(move[0] / m)*m + int(move[1] / n)] -= 1

99     # Goes one layer of minimax deeper
100    returned_score = minimax(max_depth-1, open_squares, empty_squares, m, n, not is_maximizing_player, new_score,
                                alpha, beta)[0]

102    # Changes open_squares and empty_squares back to the original state
103    open_squares.append(move)
104    empty_squares["row"][move[0]] += 1
105    empty_squares["column"][move[1]] += 1
106    empty_squares["region"][int(move[0] / m)*m + int(move[1] / n)] += 1

108    # If the score of this move going deeper is better, this becomes the best move with the best score
109    if function(returned_score, best_score):
110        best_score = returned_score
111        best_move = move

113    # Does the alpha-beta pruning
114    if is_maximizing_player: alpha = max(alpha, best_score)
115    else: beta = min(beta, best_score)
116    if alpha >= beta:
117        break

119    return best_score, best_move

121 def get_open_squares(board: SudokuBoard) -> list:
122     """
123     For the current board, gets all square that are still empty.
124     @param board: The board this should be done on.
125     @return: a list of all empty coordinates as sets.
126     """
127     open_squares = []
128     for i in range(board.N):
129         for j in range(board.N):
130             if board.get(i,j) == SudokuBoard.empty:
131                 open_squares.append((i,j))

133     return open_squares

135 def get_empty_squares(board: SudokuBoard) -> dict:
136     """
137     For the current board, gets the number of empty squares for each row/column/region.
138     @param board: The board this should be done on.
139     @return: A dictionary with keys: "row", "column" and "region"; and values being lists with the number of empty squares per group.
140     """
141     # Calculates the number of empty squares per row
142     empty_row = []
143     for i in range(board.N):
144         current_empty_row = 0
145         for j in range(board.N):
146             current_empty_row += board.get(i,j) == SudokuBoard.empty

148     empty_row.append(current_empty_row)

150     # Calculates the number of empty squares per column
151     empty_column = []
152     for i in range(board.N):
153         current_empty_column = 0
154         for j in range(board.N):
155             current_empty_column += (board.get(j,i) == SudokuBoard.empty)

157     empty_column.append(current_empty_column)

159     # Calculates the number of empty squares per region
160     empty_region = []
161     for i in range(board.N):
162         current_empty_region = 0
163         for j in range(board.N):
164             row = int(i/board.m)*board.m + int(j/board.n)
165             column = (i%board.m)*board.n + (j%board.n)
166             current_empty_region += (board.get(row,column) == SudokuBoard.empty)

168     empty_region.append(current_empty_region)

170     return {"row": empty_row, "column": empty_column, "region": empty_region}

172 def get_numbers_left(board: SudokuBoard):
173     """
174     For the current board, gets the numbers not yet in a group for each row/column/region.
175     @param board: The board this should be done on.
176     @return: A dictionary with keys: "row", "column" and "region"; and values being lists with the numbers unused per group.
177     """
178     # Calculates the missing numbers for each row
179     rows = []
180     for row in range(board.N):
181         this_row = []
182         for column in range(board.N):
183             this_row.append(board.get(row, column))
184         rows.append([x for x in range(1,board.N+1) if x not in set(filter((0).__ne__, this_row))])

186     # Calculates the missing numbers for each column
187     columns = []
188     for column in range(board.N):

```

```

189         this_column = []
190         for row in range(board.N):
191             this_column.append(board.get(row, column))
192         columns.append({x for x in range(1,board.N+1) if x not in set(filter((0).__ne__, this_column))})

194     # Calculates the missing numbers for each region
195     regions = []
196     for region in range(board.N):
197         this_region = []
198         for value in range(board.N):
199             row = int(region/board.m)*board.m + int(value/board.n)
200             column = (region%board.m)*board.n + (value%board.n)
201             this_region.append(board.get(row, column))
202         regions.append({x for x in range(1,board.N+1) if x not in set(filter((0).__ne__, this_region))})

204     return {"rows": rows, "columns": columns, "regions": regions}

206 def solve_sudoku(board, open_squares, numbers_left):
207     """
208     Iteratively gives a solution to the given sudoku.
209     First, fills in any squares where only one number is possible, then randomly guesses.
210     @param open_squares: A list containing all still open squares / possible moves.
211     @param empty_squares: A dictionary containing the missing numbers for each group.
212     @return: A filled board.
213     """
214     # Finds all squares where only one number is possible
215     result = []
216     for move in open_squares:
217         possibilities = set(numbers_left["rows"][move[0]] & numbers_left["columns"][move[1]] & \
218             numbers_left["regions"][int(move[0] / board.m)*board.m + int(move[1] / board.n)])
219         if len(possibilities) == 1:
220             number = next(iter(possibilities))

222     # If this is the case, a previous guess was wrong
223     elif len(possibilities) == 0:
224         return -1

226     # If squares can be filled in, do so and start back at the beginning
227     if result != []:
228         for i in result:
229             board.put(i[0][0], i[0][1], i[1])
230             open_squares.remove(i[0])
231             numbers_left["rows"][i[0][0]].remove(i[1])
232             numbers_left["columns"][i[0][1]].remove(i[1])
233             numbers_left["regions"][int(i[0][0] / board.m)*board.m + int(i[0][1] / board.n)].remove(i[1])

235     return solve_sudoku(board, open_squares, numbers_left)

237     # If no squares can be filled in, keep making a guess until you hit a correct one
238     elif board.empty in board.squares:
239         iterator = iter(possibilities)
240         for number in iterator:
241             new_board = deepcopy(board)
242             new_board.put(move[0], move[1], number)

244             new_open_squares = deepcopy(open_squares)
245             new_open_squares.remove(move)

247             new_numbers_left = deepcopy(numbers_left)
248             new_numbers_left["rows"][move[0]].remove(number)
249             new_numbers_left["columns"][move[1]].remove(number)
250             new_numbers_left["regions"][int(move[0] / board.m)*board.m + int(move[1] / board.n)].remove(number)
251             result = solve_sudoku(new_board, new_open_squares, new_numbers_left)

253         if result != -1:
254             return result

256     # If no possible number worked, a previous guess was wrong
257     return -1

259     # If the board is full, return
260     return board

262 # Adds three function as methods of SudokuBoard for ease of use
263 SudokuBoard.get_open_squares = get_open_squares
264 SudokuBoard.get_empty_squares = get_empty_squares
265 SudokuBoard.get_numbers_left = get_numbers_left

```
